



Summary & Leaderboard


REGIME v1 · 20 TRADING DAYS · \$100K INCEPTION

Of the four models with an estimable clean-window return, only **GPT beat SPY: +7.66%** against the benchmark's **+6.37%**.

PRIMARY LEADERBOARD · CLEAN-PILOT START (APRIL 23), DE-FABRICATED Estimable returns · integrity caveats noted

RANK	MODEL	CLEAN-WINDOW · 4/23 BASIS	VS SPY	STATUS
1	GPT-5.4 High-frequency, low-reversal directional	+7.66%	+1.29 ▲	ESTIMABLE
-	SPY passive benchmark The line every estimable model is measured against	+6.37%	-	BENCHMARK
2	DeepSeek V4 ‡ Model-identity spliced mid-window (Page 5)	+5.86%	-0.51	EXPLORATORY ‡
3	Gemini 3.1 Pro Performance shown; behavioral not characterized	+4.94%	-1.43	PERFORMANCE ONLY
4	Grok 4.20 Active two-way trader	+4.59%	-1.78	ESTIMABLE
HELD OUT · NON-ESTIMABLE, NOT OMITTED FOR RANKING				
-	Claude Sonnet 4.6 Phantom cash carried forward, corrupt book (Page 5)	-	-	NON-ESTIMABLE
-	Claude Opus 4.6 Phantom shorts carried forward, corrupt book (Page 5)	-	-	NON-ESTIMABLE

EQUITY CURVE · INDEXED TO 100 AT APRIL 9 INCEPTION Dashed = non-estimable books (Sonnet, Opus) & SPY benchmark



April 9–22 shakedown window carries fabricated launch value (equity_curve_carries_shakedown_fabrication = true). The **4/23 cut** is the leaderboard convention; everything left of it is quarantined.

FOR CONTINUITY ONLY · SINCE APRIL 9 INCEPTION · CARRIES FABRICATED LAUNCH VALUE

DeepSeek	+22.67%
GPT-5.4	+19.52%
Grok	+16.27%
Opus	+13.64%
Sonnet	+11.65%
SPY	+11.26%
Gemini	+10.00%

Not the primary measure

The since-inception order does not survive the clean cut. DeepSeek leads here but trails SPY in clean trading; Gemini, last since inception, rises to third among estimable models, its weak headline was the launch-day artifact, not clean-period trading.

Five of six models hold the majority of their since-inception return inside the corrupted window:
Sonnet 77% · Grok 74% · DeepSeek 70% · Opus 68% · GPT 57%; Gemini the exception at 36% (SPY 41%).

Performance & Risk Detail


DESCRIPTIVE SHARPE · $\pm$ 3.68% · NOT DEFLATED

In a quiet bull tape, holding **SPY** was more risk-efficient than any estimable model: its descriptive Sharpe of **6.11** sits above all four.

RISK METRICS · MAY CALENDAR MONTH


MODEL	DESCRIPTIVE SHARPE	VOLATILITY (ANN.)	MAX DD (EOD)	STATUS
SPY benchmark	6.11	–	–	BENCHMARK
GPT-5.4 · best Sharpe; risk outran, not avoided	4.65	0.175	–1.94%	ESTIMABLE
Grok 4.20 · daily risk basis-sensitive (de-fab ~3.9)	3.30	0.130	–1.92%	ESTIMABLE
Gemini 3.1 Pro · 43% coverage gap; DD understates true path	2.81	0.140	–1.88%	PERFORMANCE ONLY
DeepSeek V4 · book clean, but spans model splice	2.46	0.114	–1.34%	EXPLORATORY
Claude Sonnet 4.6 · Claude Opus 4.6: Sharpe, volatility & drawdown computed off corrupt carried-forward books	–	–	–	NON-ESTIMABLE

EQUITY CURVE RECAP · SINCE-INCEPTION, FABRICATION WINDOW SHADED



UNDERWATER · INCEPTION-ANCHORED DRAWDOWN · ESTIMABLE SET

No model approached the 30% halt; a low-stress month



RISK READ

The substantive risk-side finding is the Sharpe comparison: in a calm up-month, the passive benchmark beat every estimable model on risk-adjusted return. That is the honest frame for the month; no skill claim is asserted from ~26 clean trading days of Phase-A pilot.


BASIS CAVEATS

Grok & Gemini daily risk metrics are basis-sensitive and kept raw. DeepSeek’s Sharpe/DD are real (clean book) but span the splice. Sonnet/Opus excluded, corrupt books. All sit well below SPY’s 6.11 on any basis.

Cross-Model Behavioral & RQ_I Herding

RQ1 · DECISION CONVERGENCE

The month-one cross-model result is a **data-integrity finding, not a herding result**: three integrity gates collapse the clean cohort to a single pair.



INCLUSION GATE MATRIX

MODEL	COMPLETENESS	BOOK INTEGRITY	MODEL IDENTITY	RQ ELIGIBILITY
GPT-5.4	✓ 100%	✓ clean	✓ stable	PRIMARY
Grok 4.20	✓ 99.6%	✓ clean	✓ stable	PRIMARY
DeepSeek V4	✓ 100%	✓ clean	X splice	EXPLORATORY
Gemini 3.1 Pro	X 56.8%	✓ clean	✓ stable	EXPLORATORY
Claude Sonnet 4.6	✓ 99.6%	X corrupt	✓ stable	NON-ESTIMABLE
Claude Opus 4.6	✓ 99.2%	X corrupt	✓ stable	NON-ESTIMABLE

RQ1 HERDING · PRIMARY

Not estimable · deferred to Phase B

Cross-model herding is a property of how an ensemble co-moves; one pair cannot characterize it. The full test waits for Phase B’s fresh clean-book, single-identity cohort. That the framework caught a silently repointed model alias mid-experiment is the month-one result worth recording.

ALL-MODEL CONCORDANCE · EXPLORATORY ONLY

0.596 action concordance

Excess over chance

+0.125

Permutation p

0.002

n pairwise (bull)

647

Retained in the layer with all four exclusion reasons recorded, not promoted to a primary estimate.

EXCLUSIONS, BY REASON · KEPT DISTINCT, NOT COLLAPSED

Gemini Incomplete decision record (43% tick failure); decision-based data exploratory.	DeepSeek Model-identity splice inside the clean window; spans two underlying models.	Sonnet · Opus Corrupt carried-forward books; decision-based data non-estimable.
--	--	---

Methodology, Data Integrity & RQ Update

SOURCE COMMIT b67ac5c2

LAUNCH-WINDOW STATE CORRUPTION

Contained as flow, persistent as state for two models

A share- and cash-conservation audit of all 2,800+ decision runs isolated every state-discontinuity *event*: all are confined to the April 9–22 shakedown, with **zero events on or after April 23** for all six models. No new corruption enters the clean window, but “event-free” certifies *flow*, not a clean *book*.

FLOW · CLEAN FOR ALL SIX

Zero new fabrication events after April 23. Clean-window returns are trustworthy for all four (Gemini & Grok after a small de-fabrication correction). Risk-adjusted metrics are reliable for GPT, Grok & DeepSeek; Gemini's understate its true path on the 43% coverage gap (Page 2).

STATE · CORRUPT FOR SONNET & OPUS

Shakedown fabrication was carried forward: ~\$297k phantom cash (Sonnet), ~\$575k phantom shorts (Opus), and compounds through all of Phase A. They traded the clean window against corrupt books; performance is treated as non-estimable.

TWO BUGS · BOTH INSIDE THE SHAKEDOWN

Launch-day state clobber

4/9 → fix 4/10 · d2e862ca

Hit all six models' first trading day, the source of the implausible day-one prints (DeepSeek +9.04%, Gemini −3.75%; true ≈ +1.5%).

Cross-model state commingling

4/10–4/21 → fix 4/21 · cacd8058

Concentrated in Sonnet & Opus, the source of the carried-forward phantom state.

PER-MODEL FABRICATION · ALL INSIDE 4/9–22

MODEL	EVENTS	GROSS USD
Sonnet	28	\$233,010
Opus	32	\$187,658
DeepSeek	1	\$7,570
Gemini	1	\$5,318
Grok	1	\$4,989
GPT	1	\$1,104

THIRD INTEGRITY AXIS · MODEL IDENTITY

The provider silently repointed the deepseek-reasoner alias to a sub-frontier model; undetected ~4 weeks (month-boundary-only check).



Performance figure real (book clean): retained with a model-spliced flag, not dropped. Decision-based data spans two underlying models → moved to exploratory.

GATE-SCOPE REFINEMENT

The book-integrity gate certifies clean-window *point-return* estimability, not *daily risk-path* estimability: a small residual phantom in a volatile name leaves point returns clean while Sharpe/drawdown stay basis-sensitive (Grok, Gemini). Flagged for the September methodology pass.

RQ UPDATE · INCLUSION GATES RATIFIED · ≥80% COMPLETENESS · UNCORRUPTED BOOK · STABLE IDENTITY → MAY PASSING: GPT, GROK

RQ1 · Herding

NOT ESTIMABLE

Single-pair cohort; deferred to Phase B. All-model concordance 0.596 retained exploratory.

RQ2 · Disposition

PER-MODEL PRIMARY

GPT −0.157, Grok −0.06. Six-model pool superseded, not reported.

RQ3 · Calibration

PER-MODEL PRIMARY

GPT +0.10, Grok −0.13. Six-model pool superseded, not reported.

RQ4 · FF5 + Momentum

OPEN

Factor data ends 3/31, 0 aligned days.

RQ5 · Drawdown path

TESTING

No 10% drawdown contrast; β −0.014, p 0.118 (accumulating).

RQ6 · Reproducibility

OPEN

No determinism-probe data yet.

Risk-free rate 0.0368 (3-mo T-bill, as-of inception 4/9) · descriptive monthly Sharpe, not deflated · all figures Phase A pilot / exploratory, non-confirmatory.